

Portfolio Component: ML with SKLearn

Neo Zhao

CS 4375 - Introduction to Machine Learning

▼ 1) Read the Auto Data

```
import pandas as pd
import seaborn as sb

# 1.a.
autoData = pd.read_csv('Auto.csv', sep = ',')

# 1.b.
autoData.head()

# 1.c.
autoData.shape
```

↗ (392, 9)

▼ 2) Data Exploration

```
# 2.a.
autoData.describe()
```

`mpg` `cylinders` `displacement` `horsepower` `weight` `acceleration`

2.b.

MPG

- Range = 9 to 46.6
- Mean = 23.45

Cylinders

- Range = 3 to 8
- Mean = 5.47

Displacement

- Range = 68 to 455
- Mean = 194.41

Horsepower

- Range = 46 to 230
- Mean = 104.47

Weight

- Range = 1613 to 5140
- Mean = 2977.58

Acceleration

- Range = 8 to 24.8
- Mean = 15.55

Year

- Range = 70 to 82
- Mean = 76.01

Origin

- Range = 1 to 3
- Mean = 1.58

▼ 3) Explore Data Types

3.a.

```
print("Data Types before Modification")  
print(autoData.dtypes)
```

```
# 3.b.
autoData.cylinders = autoData.cylinders.astype('category').cat.codes

# 3.c.
autoData.origin = autoData.origin.astype('category')

# 3.d.
print("\nData Types after Modification")
print(autoData.dtypes)
```

```
Data Types before Modification
cylinders          int8
displacement      float64
horsepower        int64
weight            int64
acceleration      float64
year              float64
origin            category
mpg_high          category
dtype: object
```

```
Data Types after Modification
cylinders          int8
displacement      float64
horsepower        int64
weight            int64
acceleration      float64
year              float64
origin            category
mpg_high          category
dtype: object
```

▼ 4) Deal with NAs

```
# 4.a.
autoData = autoData.dropna()

# 4.b.
autoData.shape

(389, 9)
```

▼ 5) Modify Columns

```
# 5.a.
```

```

avg = autoData['mpg'].mean()
autoDataNew = autoData.loc[autoData.mpg > 1].copy()

autoDataNew.loc[:, 'mpg_high'] = [0 if x < avg else 1 for x in autoDataNew['mpg']]
autoData = autoDataNew
autoData.mpg_high = autoData.mpg_high.astype('category')

# 5.b.
autoData = autoData.drop (columns = ['mpg', 'name'])

# 5.c.
print(autoData.head())

print(autoDataNew.head())

```

	cylinders	displacement	horsepower	weight	acceleration	year	origin \
0	4	307.0	130	3504	12.0	70.0	1
1	4	350.0	165	3693	11.5	70.0	1
2	4	318.0	150	3436	11.0	70.0	1
3	4	304.0	150	3433	12.0	70.0	1
6	4	454.0	220	4354	9.0	70.0	1

	mpg_high
0	0
1	0
2	0
3	0
6	0

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
0	18.0	4	307.0	130	3504	12.0	70.0	
1	15.0	4	350.0	165	3693	11.5	70.0	
2	18.0	4	318.0	150	3436	11.0	70.0	
3	16.0	4	304.0	150	3433	12.0	70.0	
6	14.0	4	454.0	220	4354	9.0	70.0	

	origin	name	mpg_high
0	1	chevrolet chevelle malibu	0
1	1	buick skylark 320	0
2	1	plymouth satellite	0
3	1	amc rebel sst	0
6	1	chevrolet impala	0


▼ 6) Data Exploration (Graphs)

```

# 6.a.
sb.catplot(x = 'mpg_high', kind = 'count', data = autoData)

```


```
<seaborn.axisgrid.FacetGrid at 0x7f6292cebb90>
```



```
# 6.b.
```

```
sb.relplot(x = 'horsepower', y = 'weight', data = autoData, hue = autoData.mpg_high)
```


```
<seaborn.axisgrid.FacetGrid at 0x7f6292850990>
```



```
# 6.c.
```

```
sb.boxplot(x = 'mpg_high', y = 'weight', data = autoData)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f62927abed0>



6.d.

6.a.

- There are more samples of "low_mpg" than "high_mpg"; however, they are almost evenly split

6.b.

- A lighter car has less horsepower than any heavier car. Therefore, heavier cars are likely to consume gas faster than lighter cars.

#6.c.

- "high_mpg" seems to have several outliers while "low_mpg" doesn't have any.

▼ 7) Train/Test Split

7.a. & 7.b. & 7.c. & 7.d.

```
from sklearn.model_selection import train_test_split
```

```
X = autoData.loc[:, autoData.columns != 'mpg_high']
```

```
y = autoData.mpg_high
```

```
# Split into 80/20 with seed = 1234
```

```
xTrain, xTest, yTrain, yTest = train_test_split(X, y, test_size = 0.2, random_state = 1234)
```

```
print('Train Size:', xTrain.shape)
```

```
print('Test Size:', xTest.shape)
```

```
Train Size: (311, 7)
```

```
Test Size: (78, 7)
```

▼ 8) Logistic Regression

8.a.

```

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report

clf = LogisticRegression(solver = 'lbfgs', max_iter = 300)
clf.fit(xTrain, yTrain)
clf.score (xTrain, yTrain)

# 8.b.
pred = clf.predict(xTest)

# 8.c.
# 0 = "low_mpg"
# 1 = "high_mpg"

print(classification_report(yTest, pred))

```

	precision	recall	f1-score	support
0	1.00	0.84	0.91	50
1	0.78	1.00	0.88	28
accuracy			0.90	78
macro avg	0.89	0.92	0.89	78
weighted avg	0.92	0.90	0.90	78

▼ 9) Decision Trees

```

# 9.a.
from sklearn.tree import DecisionTreeClassifier
from sklearn.datasets import load_iris
from sklearn import tree
from matplotlib import pyplot as plt

clf2 = DecisionTreeClassifier()
clf2.fit(xTrain, yTrain)
clf2.score(xTrain, yTrain)


# 9.b.
pred2 = clf2.predict(xTest)

# 9.c.
print(classification_report(yTest, pred))

# 9.d.
plt.figure(figsize = (10, 10))
tree.plot_tree(clf2)
plt.show()

```

	precision	recall	f1-score	support
0	1.00	0.84	0.91	50
1	0.78	1.00	0.88	28
accuracy			0.90	78
macro avg	0.89	0.92	0.89	78
weighted avg	0.92	0.90	0.90	78



10) Neural Network

10.a.

```
from sklearn import preprocessing
from sklearn.neural_network import MLPClassifier

scaler = preprocessing.StandardScaler().fit(xTrain)

xTrainScale = scaler.transform(xTrain)
```



```

xTestScale = scaler.transform(xTest)

clf3 = MLPClassifier (solver = 'lbfgs', hidden_layer_sizes = (6), max_iter = 1500, random_state=42)
clf3.fit(xTrainScale, yTrain)
clf3.score (xTrainScale, yTrain)

# 10.b.
pred3 = clf3.predict(xTestScale)

print(classification_report(yTest, pred3))

# 10.c.
clf4 = MLPClassifier(solver = 'sgd', hidden_layer_sizes = (3,), max_iter = 1500, random_state=42)
clf4.fit(xTrainScale, yTrain)
clf4.score (xTrainScale, yTrain)

# 10.d.
pred4 = clf4.predict(xTestScale)

print(classification_report(yTest, pred4))

```

	precision	recall	f1-score	support
0	0.94	0.90	0.92	50
1	0.83	0.89	0.86	28
accuracy			0.90	78
macro avg	0.89	0.90	0.89	78
weighted avg	0.90	0.90	0.90	78

	precision	recall	f1-score	support
0	0.93	0.80	0.86	50
1	0.71	0.89	0.79	28
accuracy			0.83	78
macro avg	0.82	0.85	0.83	78
weighted avg	0.85	0.83	0.84	78

10.e.

- The first Neural Network performed better by about 7% compared to the second one. I used to same amount of iterations for both networks, but the hidden layers are different for each of the networks. Logistic Regression and Decision Tree performed the same, which might mean there's a mistake somewhere.

11) Analysis

11.a. Which algorithm performed better?

- I think the Neural Networks performed the best

11.b. Compare Accuracy, Recall, and Precision metrics by class

- Logistic Regression
 - Accuracy = 90%
 - "low_mpg" Recall = 84%
 - "high_mpg" Recall = 100%
 - "low_mpg" Precision = 100%
 - "high_mpg" Precision = 78%
- Decision Tree
 - Accuracy = 90%
 - "low_mpg" Recall = 84%
 - "high_mpg" Recall = 100%
 - "low_mpg" Precision = 100%
 - "high_mpg" Precision = 78%
- Neural Network #1
 - Accuracy = 90%
 - "low_mpg" Recall = 90%
 - "high_mpg" Recall = 89%
 - "low_mpg" Precision = 94%
 - "high_mpg" Precision = 83%
- Neural Network #2
 - Accuracy = 83%
 - "low_mpg" Recall = 80%
 - "high_mpg" Recall = 89%
 - "low_mpg" Precision = 93%
 - "high_mpg" Precision = 71%

#11.c. Give your analysis of why the better-performing algorithm might have outperformed the other

- I believe Neural Networks perform more efficiently when working with complex data, allowing the model to outperform Logistic Regression as well as Decision Trees.

#11.d. Write a couple of sentences comparing your experiences using R versus sklearn. Feel free to express strong preferences.

- I definitely prefer R at this point, since I have worked with it more in the past and I am more comfortable with this language; however, after this assignment, I am ready to dive deeper into learning Machine Learning with Python as it is always useful to know more.

[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 4:34 PM

